

Delivered tables — preview

Not a deliverable. Layout here is arbitrary; the paper's floats are D's call. Generated by `python -m analysis.preview`.

table1_dataset

Statistic	Development	Test
Paragraphs	2000	2000
duplicated across splits	1	1
Source reports (PDFs)	49	49
shared across splits	49	49
Companies	49	49
spanning >1 report	0	n/a
Legal states observed	15 / 17	n/a
Gold rows violating the hierarchy	0	n/a
<i>Rarest classes</i>		
within_2_years	34	n/a
Misleading	2	n/a

Dataset and split statistics, regenerated from the versioned data by `analysis/audit.py`. Misleading occurs twice in the whole development set, in two different reports, and is absent from the Calibration partition of 18 of the 30 rotations. The four fields admit 120 combinations, of which the hierarchy leaves 17 legal, and no development row violates it. All 49 development reports also appear in the test split, and one paragraph text is duplicated across the two. No company contributes more than one report, so a document-disjoint split is also company-disjoint. The competition test split ships no labels, so its label-derived cells are marked n/a.

table2_main

ID	Calibration	Decoding	Weighted F1	PS	VT	ES	EQ	Tuple Acc.	Invalid %
M0	None	Independent	0.572±0.003	0.796	0.464	0.650	0.424	0.359	12.6
M1	None	Projection	0.571±0.003	0.796	0.467	0.651	0.418	0.394	0.0
M2	Global	Projection	0.574±0.004	0.796	0.493	0.650	0.418	0.428	0.0
M3	Conditional	Projection	0.574±0.005	0.796	0.501	0.651	0.412	0.431	0.0
M4	None	17-state	0.571±0.005	0.799	0.468	0.648	0.420	0.388	0.0
M5	Global	17-state	0.576±0.001	0.796	0.493	0.649	0.423	0.430	0.0
M6	Conditional	17-state	0.567±0.009	0.784	0.494	0.636	0.414	0.426	0.0

Controlled comparison of decision rules on identical base probabilities and identical test rows. Each cell is the mean over three seeds of a single score computed on all 2,000 concatenated test rows; $\pm$ is the sample standard deviation across seeds, which reflects the variability of the whole pipeline – fold assignment and training together – rather than model stability. Paired contrasts between these rows, under each metric and with their Holm-corrected p-values, are reported in Table 4. Under that variant M0 scores 0.567 against 0.572 on the official metric, while M1–M6 score identically under both: only the unconstrained arm predicts fields its own ancestors do not support.

table3_regimes

Method	Same-document	Document-disjoint	Δ	95% CI
M0	0.585	0.572	0.012	[0.004, 0.022]
M1	0.581	0.571	0.010	[0.001, 0.019]
M2	0.587	0.574	0.012	[0.001, 0.024]
M3	0.588	0.574	0.014	[0.003, 0.025]
M4	0.586	0.571	0.014	[0.005, 0.024]
M5	0.590	0.576	0.015	[0.004, 0.027]
M6	0.587	0.567	0.020	[0.011, 0.031]

Same-document versus document-disjoint evaluation. The left column measures seen-report, unseen-paragraph generalisation and matches the competition distribution, since test and development draw on the same 49 reports; the right column measures generalisation to entirely unseen reports. Δ is the gap between two estimation targets and is not a bias estimate. All seven decision rules are reported rather than a best-scoring subset: every Δ is positive and every interval excludes zero, so the gap is a property of the split and not of any one rule.

table4_contrasts

Metric	Contrast	Δ	95% CI	p_{Holm}
Weighted macro-F1 (official)	M1-M0	-0.001	[-0.006, 0.003]	1.000
	M3-M2	-0.001	[-0.006, 0.004]	1.000
	M4-M1	+0.000	[-0.005, 0.005]	1.000
	M6-M3	-0.007	[-0.017, 0.003]	0.643
	M6-M5	-0.009	[-0.017, -0.001]	0.135
Tuple accuracy	M1-M0	+0.035	[0.028, 0.043]	0.001
	M3-M2	+0.002	[-0.003, 0.007]	0.973
	M4-M1	-0.006	[-0.010, -0.002]	0.028
	M6-M3	-0.005	[-0.014, 0.004]	0.973
	M6-M5	-0.004	[-0.012, 0.004]	0.973

The five pre-specified contrasts of the analysis plan, under the two metrics the plan named before any result existed: the competition’s own weighted macro-F1 and tuple exact-match. Paired PDF-cluster bootstrap over 10,000 resamples on the document-disjoint protocol; the resampling unit is the source report, not the paragraph. Each metric is Holm-corrected as its own family of five rather than pooled: the contrasts were specified once, not once per metric. **Bold p_{Holm} marks a contrast that survives the correction; the bracketed intervals are uncorrected 95% percentile intervals and decide nothing on their own.** Two further structure-aware metrics were examined after the primary analysis and are reported as exploratory in the text, not as Holm families here. Of the five contrasts, none on Weighted macro-F1 (official) and two on Tuple accuracy survive the correction at $\alpha=0.05$. The official metric and tuple accuracy were named in the analysis plan in advance; the path-constrained variant and the hierarchical F1 were adopted after the primary analysis and are not pre-specified.

table5_metrics

ID	Weighted F1	C-wF1	hP	hR	hF	Tuple Acc.
M0	0.572	0.567	0.688	0.703	0.695	0.359
M1	0.571	0.571	0.707	0.690	0.698	0.394
M2	0.574	0.574	0.722	0.719	0.721	0.428
M3	0.574	0.574	0.723	0.722	0.722	0.431
M4	0.571	0.571	0.700	0.703	0.702	0.388
M5	0.576	0.576	0.721	0.730	0.725	0.430
M6	0.567	0.567	0.713	0.745	0.729	0.426

The same seven decision rules scored under each metric. C-wF1 is the official metric with ancestor-unsupported fields counted as false predictions; hP, hR and hF are the ancestor-based hierarchical precision, recall and F1, which are micro-averaged over path nodes rather than macro-averaged over classes. C-wF1 equals the official score for every rule whose output is legal by construction and falls below it only for the unconstrained one. The metrics disagree about the winner: M5 on the official metric and M6 on hF, which is the lowest-scoring rule on the official metric. A ranking is only meaningful stated together with the metric that produced it.

table7_resolution

Contrast	$ \Delta $	CI half-width	Interval excludes 0	p_{Holm}	Survives
M4-M1	0.000	0.005	no	1.000	no
M3-M2	0.001	0.005	no	1.000	no
M1-M0	0.001	0.005	no	1.000	no
M6-M3	0.007	0.010	no	0.643	no
M6-M5	0.009	0.008	yes	0.135	no

What this benchmark can resolve, on the metric it ranks by. Contrasts are ordered by effect size. A percentile interval that excludes zero is not a verdict: with five contrasts corrected together, the largest effect here clears zero and still does not survive. That magnitude bounds what a leaderboard difference on this benchmark can mean, and it should be read against the full range across the seven decision rules, which is 0.0088.


table6_legality_cost

Backbone	λ	M0 invalid	Enforce legality (M1–M0)		Joint decoding (M4–M1)	
			Official	Whole-row	Official	Whole-row
RoBERTa-large	0	12.55%	-0.001	+0.035	+0.000	-0.006
DeBERTa-v2-320M	0	19.75%	-0.008	+0.050	+0.013	-0.003
ELECTRA-large	0	10.20%	-0.005	+0.024	+0.007	-0.000
RoBERTa-base	0	11.77%	-0.007	+0.030	+0.008	-0.003
RoBERTa-large	0.3	5.18%	-0.001	+0.014	-0.001	-0.001
DeBERTa-v2-320M	0.3	6.53%	+0.001	+0.015	-0.002	-0.003
ELECTRA-large	0.3	3.68%	-0.000	+0.009	-0.000	-0.002

Two decision rules over one set of probabilities, on the pdf_group protocol. M0 takes each field’s argmax independently; M1 projects onto the 17 legal states top-down; M4 scores all 17 and takes the best. **M1 and M4 both emit an invalid tuple on 0% of rows in every arm** – their output space is the legal set – so the invalid column reports M0 only, where it ranges from 3.68% to 19.75%. Means over 3 seeds; paired PDF-cluster bootstrap, 10,000 resamples, seed 20260814, one resample shared within each contrast. Enforcing legality raises whole-row accuracy in all 7 arms and lowers the official score in all 4 arms trained without the structural objective; joint decoding reverses both signs. The 3 structurally trained arms show neither effect. **Exploratory.** These contrasts were named after the primary analysis and form no Holm family; bold marks an interval excluding zero, not a corrected verdict. The claim is the sign pattern across arms, not any single cell. See docs/inference_families.md.

figure1_hierarchy

(a) Task hierarchy



$1 + 4 \times (1 + 3) = 17$ legal tuples

of $2 \times 5 \times 3 \times 4 = 120$ combinations

the other 103 are hierarchy-invalid

(b) Decision routes (alternatives, not stages)

